

A Journey into Hexagon

Dissecting Qualcomm Basebands


Seamus Burke

Agenda

- [About me](#)
- Why Basebands?
- History
- Modern SoC Architecture
- Hexagon
- Cellular Stack architecture
- Analysis

About Me

- Been working in infosec for 4 years
- Interested in kernel internals, exploit development, and embedded systems
- Plays a lot of CTFs



Goals

- Understand how it works and how it interacts with Android
- Find bugs in the baseband
- Have fun

Prior Work

“Reverse Engineering a Qualcomm Baseband” - Guillaume Delugré (CCC 2011)

“Baseband Exploitation in 2013” - Ralph-Philipp Weinmann (PacSec 2013)

“Exploring Qualcomm Baseband via Modkit” - Peter Pi, XiLing Gong, GmXP (CSW 2018)

Agenda

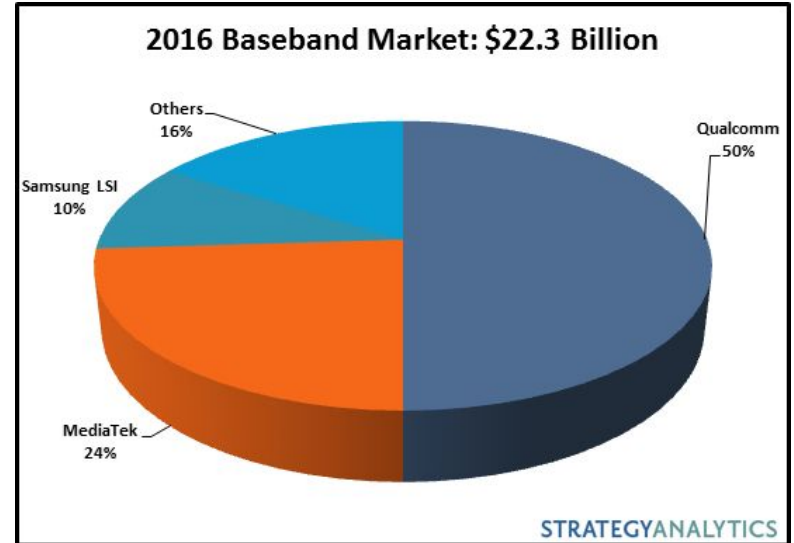

- About me
- [Why Basebands?](#)
- History
- Modern SoC Architecture
- Hexagon
- Cellular Stack architecture
- Analysis

What is a baseband exactly?

- The chip in a phone which communicates with the cellular network
- Handles the radio signal processing
- Has to support a large number of standards -
(GSM,UMTS,CDMA2k,cdmaOne, GPRS,EDGE, LTE, etc)
- The phones main external interface to the rest of the world

Why Qualcomm?

- By far the largest market share of any baseband processor
- Most high-end phones carry a Qualcomm chip



Agenda

- About me
- Why Basebands?
- History
- Modern SoC Architecture
- Hexagon
- Cellular Stack architecture
- Analysis

RTOS

- Real-time, embedded operating systems.
- 2 major categories
 - Unprotected - logical and physical addresses are the same
 - Protected - virtual memory
- Time bound, with well defined time constraints



REX - Real-time Executive

- The original kernel which ran the modem
- Designed for the 80186, then ported to ARM
- Single process, multiple threads
- Everything runs in supervisor mode



REX Quirks

- Tasks are basically the threads of REX.
- Every task has a TCB storing the the SP, priority, stack_limit, etc
- Each task has it's own stack, when tasks are suspended, context is saved in a special context frame on top of it's stack
- Pointer to the saved context stored in the tasks TCB
- TCBs stored in a doubly linked list, trivial to mess with

Moving Into the mid-2000s....

- Why did Qualcomm switch from REX?
- Well, it had it's issues:

“Programmers should use these functions only according to the interface specifications, since REX does not check any parameter for validity”

- Flat address space, lack of memory protections
 - Did they switch for security? Hah, no, debugging millions of lines of C with no memory protections was a nightmare

L4 + Iguana

- Multiple-process, multi-threaded
- Only the kernel runs in supervisor mode, everything else in userland
- A REX emulation layer is supported
 - REX tasks are L4 threads
 - No changes to the REX api, it's converted transparently
 - AMSS runs in user mode
 - Interrupts split between kernel and user mode

QuRT

- Qualcomm Real-time OS
- Used to be named BLAST, name changed as part of OpenDSP initiative
- Most of the APIs are backwards compatible, with the exception of some threading things.
- Provides pretty standard OS primitives
 - Mutexes
 - Semaphores
 - Priority Based Scheduling
- Handles the Virtual-Physical memory mapping

Mitigations? Sorta

- Complete lack of ASLR
- There is a form of DEP, can't write to code, can't execute data
- XORd stack cookies
- Heap protection
 - Different magic values for the headers of in-use and free'd blocks
- Lots of fixed addresses everywhere. The RTOS loads at the same spot every time, as does just about everything else.
- Hardcoded addresses prevalent in the code

AMSS

- Advanced Mobile Subscriber Software, drivers and protocol stacks which make up the modem
- Configured differently for different chipsets
 - Which air interface protocols are supported
 - Hardware specific drivers
 - Multimedia software
- > 60 tasks running
 - Diag, Dog, Sleep, CM, PS, DS, etc

Diagnostics

- DIAG, or Diagnostic Services provides a server that clients can request info from about AMSS
- DIAG is a REX task, usually handles requests from Serial or USB
- Packet based protocol
- Lots of useful stuff
 - Debug messages
 - OTA logs
 - Status
 - Statistics

Agenda

- About me
- Why Basebands?
- History
- Modern SoC Architecture
- Hexagon
- Cellular Stack architecture
- Analysis

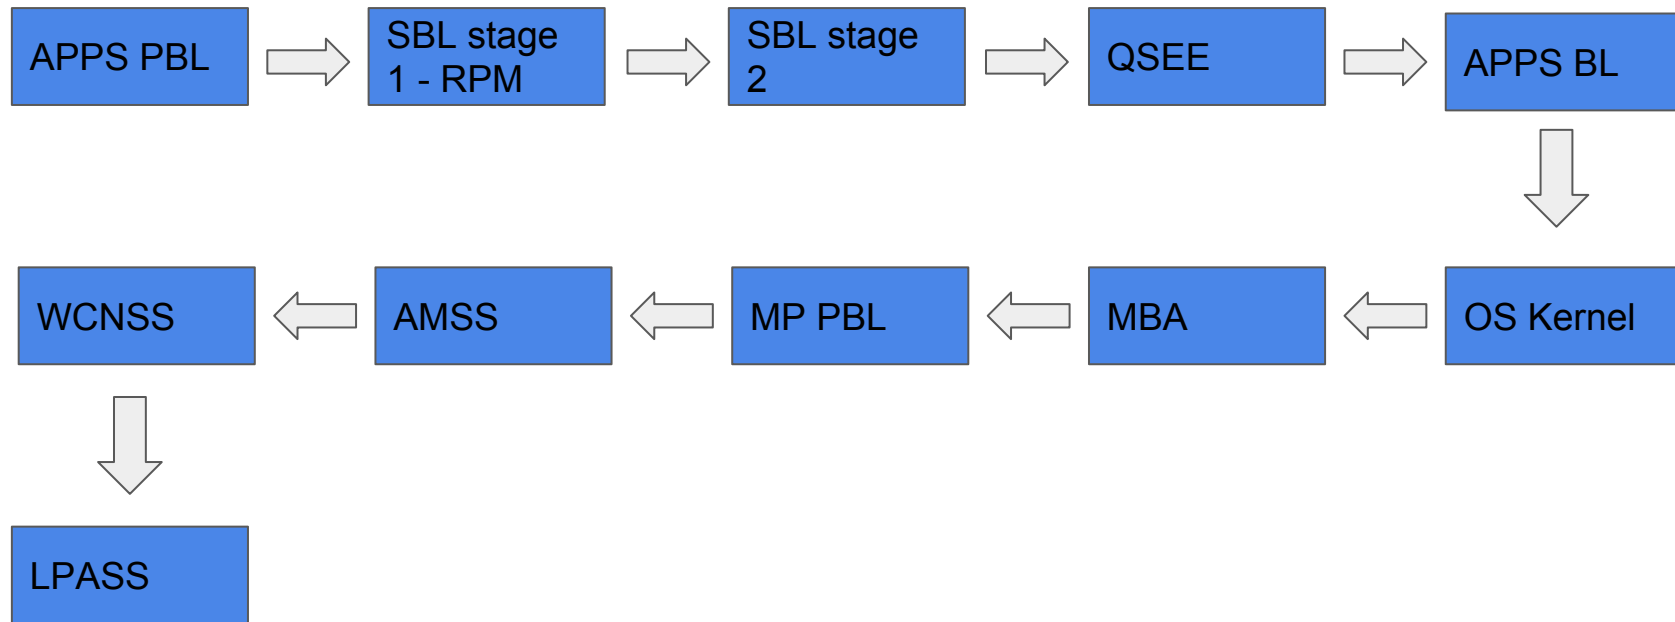
Qfuses

- Internal bank of one time programmable fuses, the QFPROM
- Publicly undocumented
- Inter-chip configuration settings, cryptographic keys
- Secure boot and TrustZone both make heavy use of these
- Hardware debugging usually disabled in prod by blowing a fuse

SoC Architecture

- Multiple interconnected subsystems
 - MPSS - Modem Processor
 - APPS - Application processor
 - RPM - Resource and Power Management
 - WCNSS - Wireless Connectivity
 - LPASS - Low Power Audio
- Baseband not the master over the application processor

SoC Boot Order



Shared Memory

- Main idea is for the Modem to write some data, and the AP pick it up
- Common APIs on linux and the modem (and other subsystems)
 - `Smem_init`, `smem_alloc`, `smem_find`, `smem_get_entry`
- SMD - Shared Memory Driver
 - Wrapper around SMEM
 - Abstracts into things like pipes
 - Separate channels for DS, DIAG, RPC, GPS

QMI

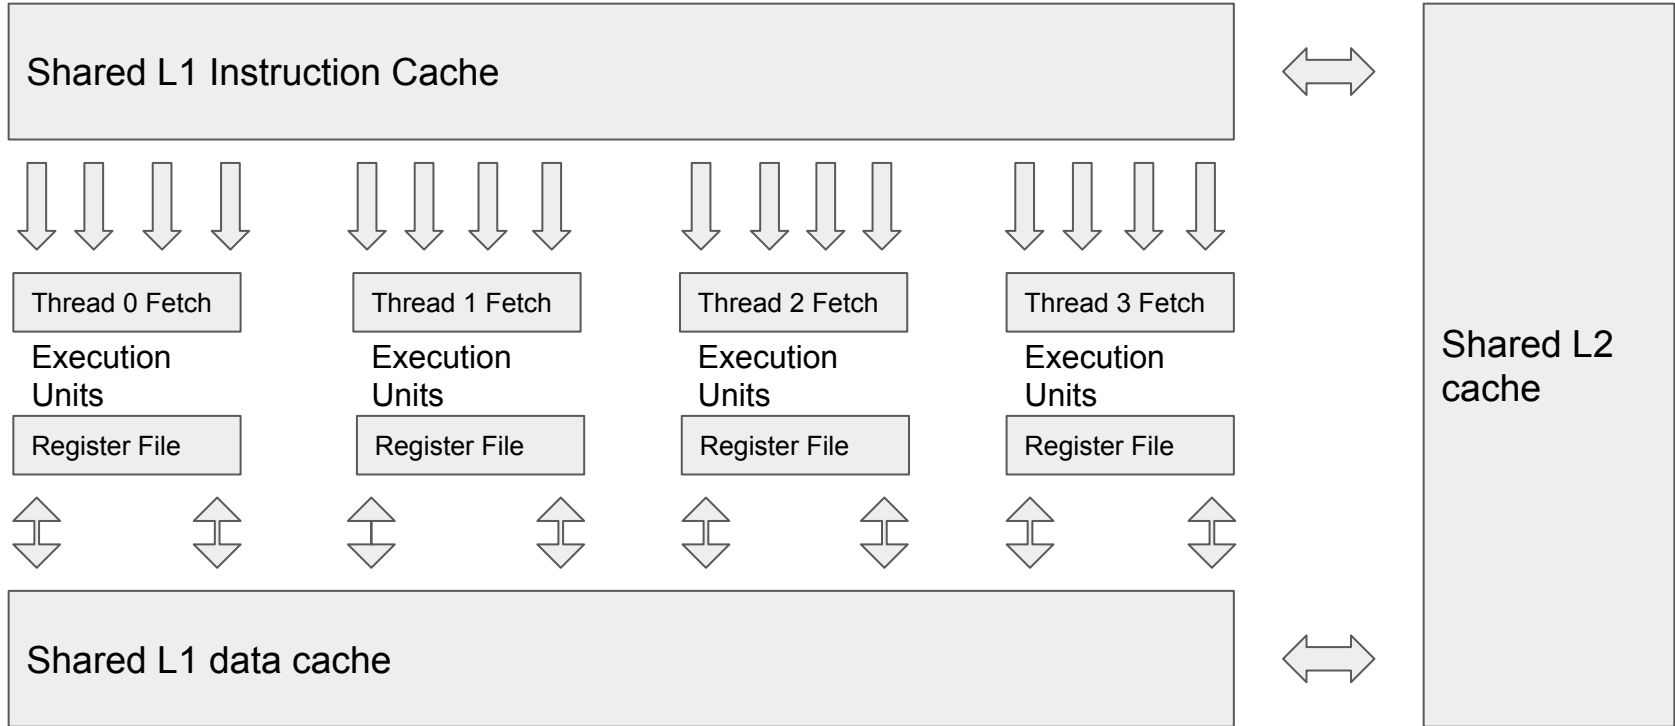
- Qualcomm MSM Interface - designed to supplant the AT cmd set
- High level interface over older protocol (DMSS)
- Used to interface with modem components, but not drive hw
- Client-server model
- Packet structure with a header and then TLV payloads

Agenda

- About me
- Why Basebands?
- History
- Modern SoC Architecture
- Hexagon
- Cellular Stack architecture
- Analysis

Hexagon

- Qualcomm's in house DSP.
- General purpose DSP. 2 on SoC (1 each for application and cellular processor)
- Separate L1 code and data caches, unified L2 cache
- Hardware threads share caches
- Instructions grouped into packets of 1-4 instructions for parallel execution



Architecture Info I

- Thirty two 32-bit GPRs
- Calling convention - R0-R5 are used for arguments
- Return values in R0
- Caller saved are R6-R15
- Callee saved R16-R27


The stack on QDSP

- R29 is Stack Pointer, R30 is Frame Pointer, R31 is Link Register
- The stack grows downwards from high to low
- Needs to be 8 byte aligned
- Several stack specific instructions
 - Allocframe - pushes LR and FP to stack, and subtracts from SP to make room for the new frames locals.
 - Deallocframe - Loads FP and LR, then fixes up SP
 - Dealloc_return - does a deallocframe and then returns to LR

Higher Address



Lower Address



Architecture Info II

- SSR - holds a variety of useful debugging info
 - ASID
 - CAUSE
 - Which BadVA
- BADVA
 - BADVA0 - exception addresses generated from slot0
 - BADVA1 - addresses generated from slot1
- ELR - holds PC value when an exception occurs

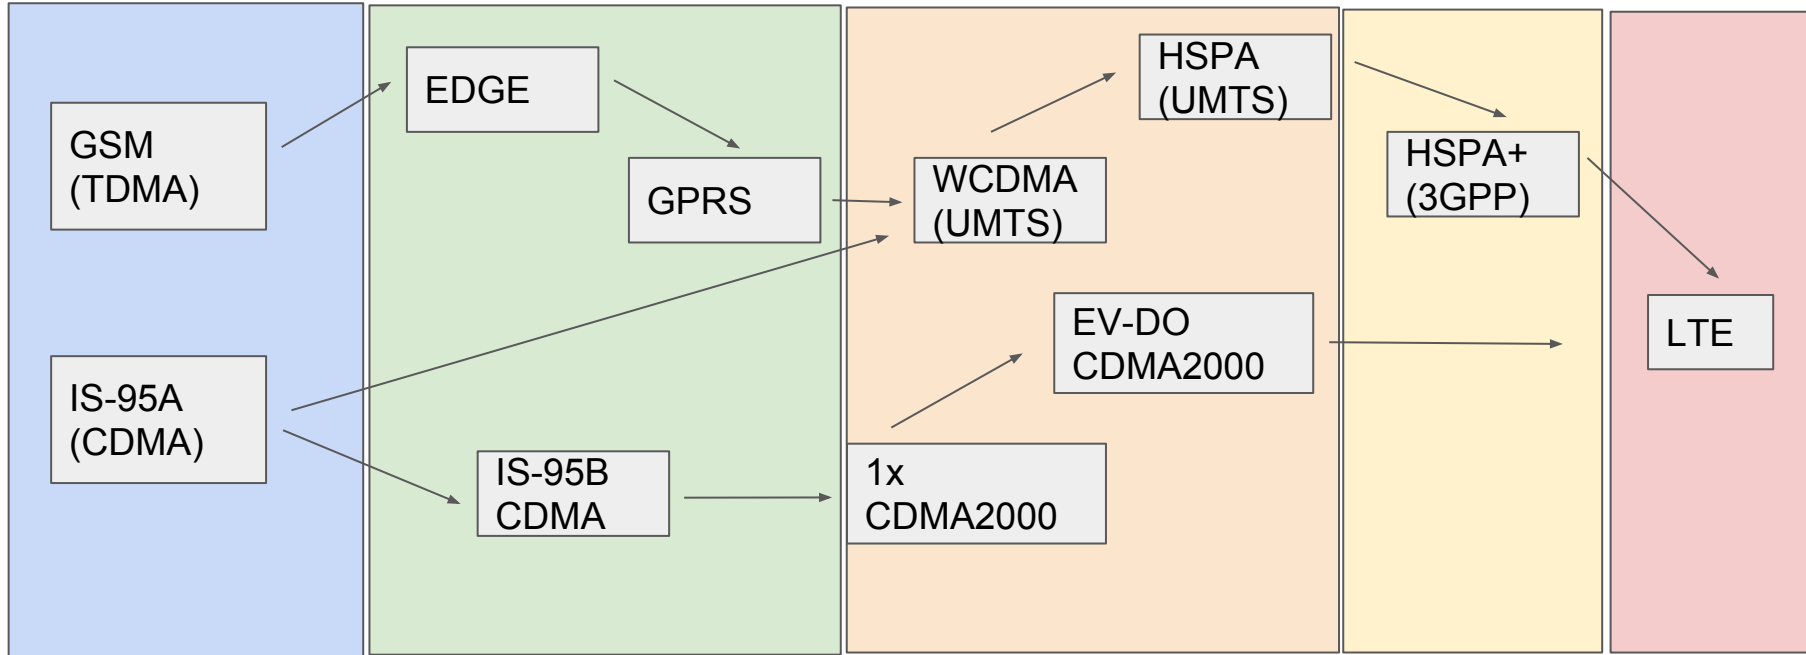
Privilege Modes and Protection Domains

- 3 Main modes
 - Kernel Mode - Access to all memory, smallest memory footprint
 - Guest OS - Access to it's own memory, and of the User segment, lots of Qcom drivers run here
 - User - Only has access to itself.
- Stack checks only done in user and guest
- Protection Domains implement separate address spaces
 - Memory mapping is (Address space ID + virtual address)
 - ASID0 is Kernel and Guest

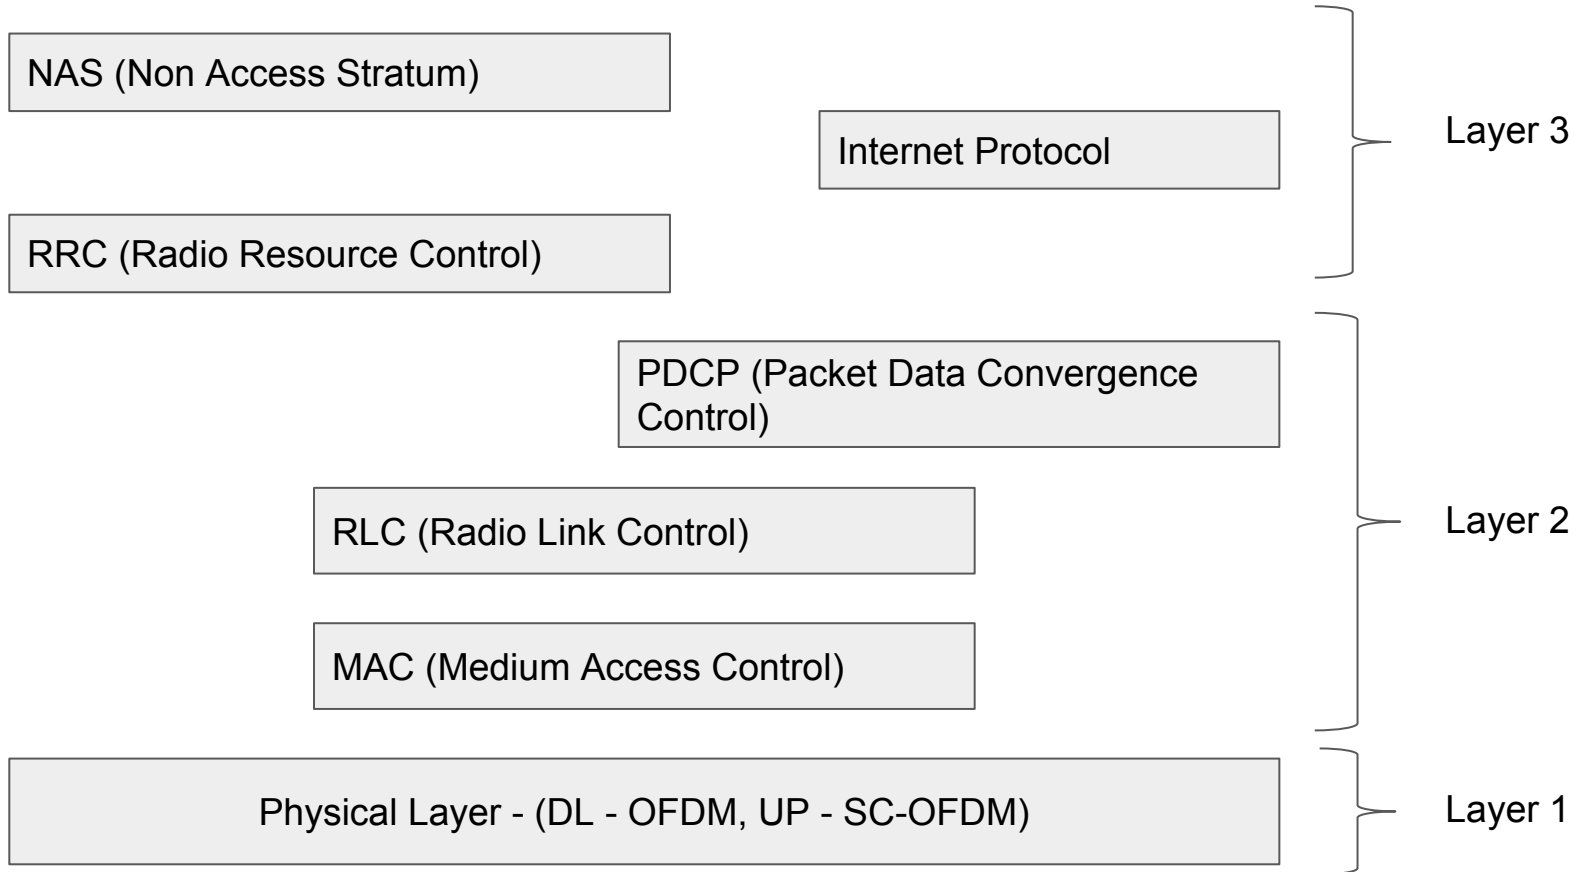
Agenda

- About me
- Why Basebands?
- History
- Modern SoC Architecture
- Hexagon
- Cellular Stack architecture
- Analysis

Evolution of Cellular Standards

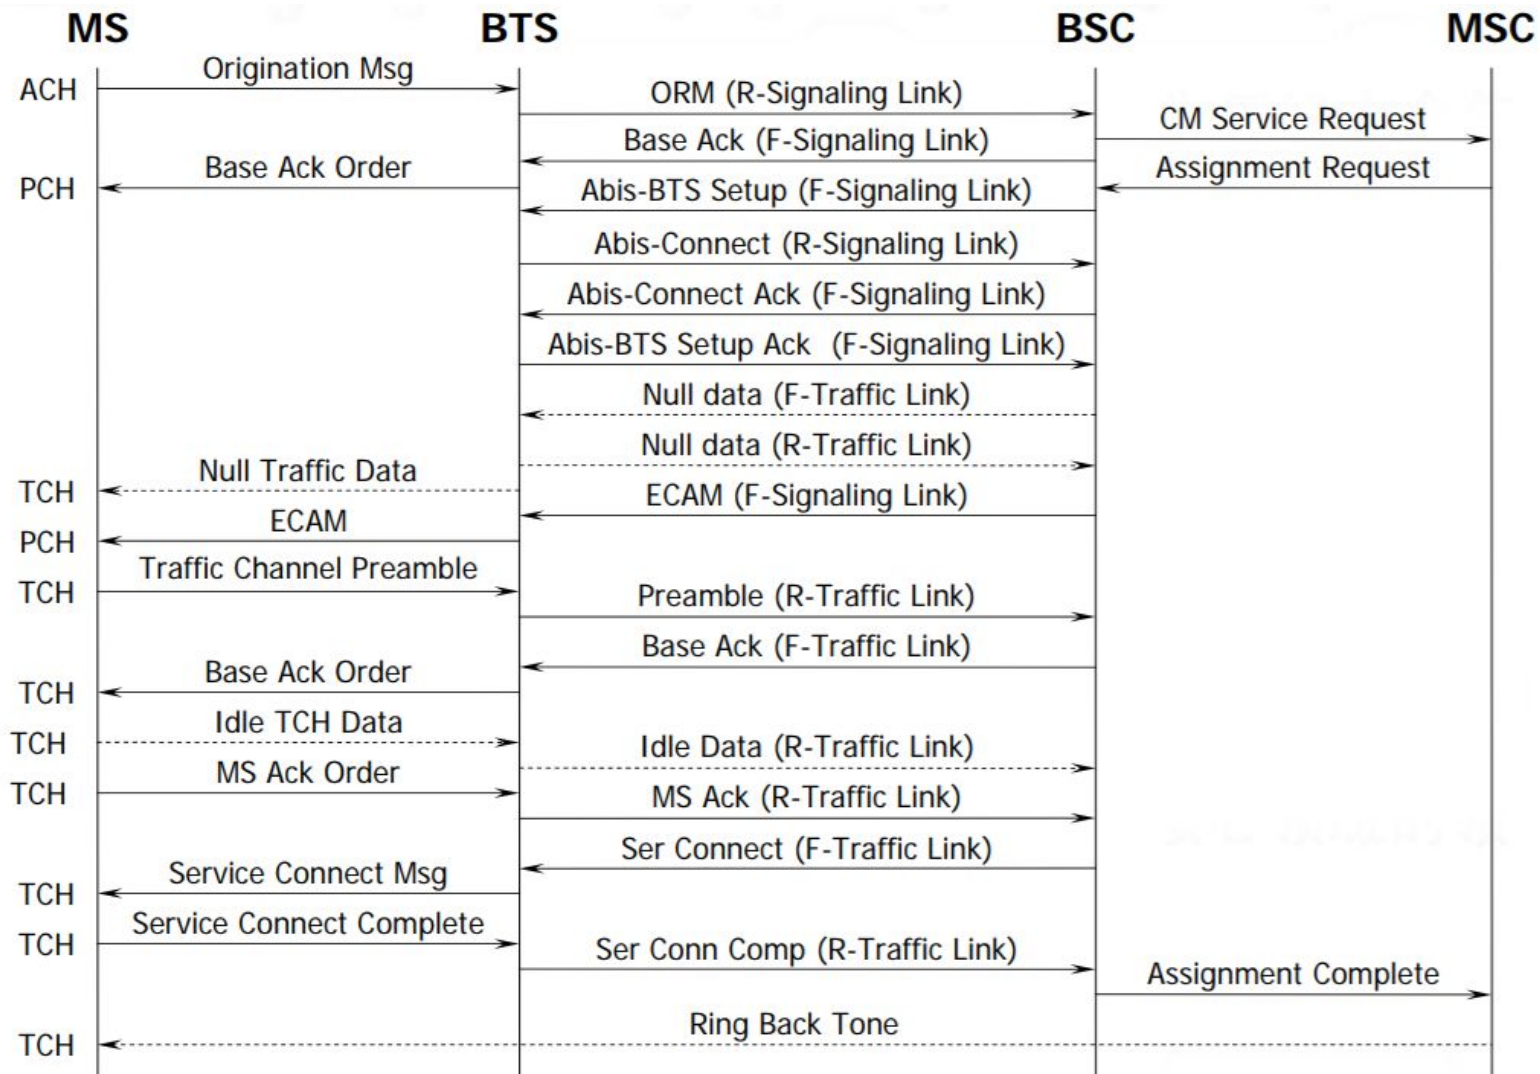


Cellular Protocol Layers



Call flow

- There are a dozen different ways the cell stack can make/receive a call
 - 1x voice call
 - 1x data call
 - HDR call
 - GSM voice call
 - GPRS data call
 - WCDMA voice call
 - WCDMA data call
 - TD-SCDMA call
 - LTE data call
- Multiple ways to do the same thing implies added complexity, and these aren't as simple as a 3-way handshake to begin with



RTFS

- Best place to start is the standards
- Don't specify implementation, and there are lots of features to implement
- 3GPP is the governing body
 - Composed of 7 telecom orgs (ETSI, ARIB, ATIS, CCSA, TSDSI, TTA, TTC)
- The standards are freely available on the web
- (Very long)

Digging Into the Standards for Bugs

- Plenty of LV and TLV options everywhere. Good place to start
- Can be tricky to find out how to trigger them

8	7	6	5	4	3	2	1	
User-user IE								octet 1
Length of user-user contents								octet 2
User-user protocol discriminator								octet 3
User-user information								octet 4*
								octet N*

User-user protocol discriminator (octet 3)								
8	7	6	5	4	3	2	1	Bits
0	0	0	0	0	0	0	0	User specific protocol (Note 1)
0	0	0	0	0	0	0	1	OSI high layer protocols
0	0	0	0	0	0	1	0	X.244 (Note 2)
0	0	0	0	0	0	1	1	Reserved for system management convergence function
0	0	0	0	0	1	0	0	IA5 characters (Note 3)
0	0	0	0	0	1	1	1	rate adaption according to ITU-T Rec. V.120 [61]

	Mobile identity	Mobile identity 10.5.1.4	M	LV	6-9
	Old routing area identification	Routing area identification 10.5.5.15	M	V	6
	MS Radio Access capability	MS Radio Access capability 10.5.5.12a	M	LV	6-51
3	Old P-TMSI signature	P-TMSI signature 10.5.5.8	O	TV	4
7	Requested READY timer value	GPRS Timer 10.5.7.3	O	TV	2
3-7	TMSI status	TMSI status 10.5.5.4	O	TV	1
33	PS LCS Capability	PS LCS Capability 10.5.5.22	O	TLV	3-4
	Mobile station classmark 2	Mobile station classmark 2 10.5.1.6	O	TLV	5
	Mobile station classmark 3	Mobile station classmark 3 10.5.1.7	O	TLV	2-34
	Supported Codecs	Supported Codec List 10.5.4.32	O	TLV	5-n
	UE network capability	UE network capability 10.5.5.26	O	TLV	4-15
	Additional mobile identity	Mobile identity 10.5.1.4	O	TLV	7

Agenda

- About me
- Why Basebands?
- History
- Modern SoC Architecture
- Hexagon
- Cellular Stack architecture
- Analysis

Disassembly options

- There are several options out there for disassembling hexagon code
 - <https://github.com/gsmk/hexagon>
 - <https://github.com/programa-stic/hexag00n>
 - <https://github.com/rpw/hexagon>
 - Qcom provided patches for GNU binutils
 - Llvm, codebench, etc
- I found the GSMK plugin the fastest to setup and get running
- I have a very rough binary ninja based disassembler I wrote

Analysis

- Qdsp6sw.mbn - Holds the modem firmware and QuRT
- Not small -

```
seamus@RIL:~/Desktop/modem$ ls -al
total 29356
drwxrwxr-x  2 seamus seamus   4096 Jul 11 00:43 .
drwxr-xr-x 14 seamus seamus   4096 Jul 11 00:43 ..
-rw-rw-r--  1 seamus seamus 30050091 Jul 11 00:43 qdsp6sw.mbn
```

- It has tens of thousands of functions to sort through
- Where to start?

```
The initial autoanalysis has been finished.
The total number of functions is 86217
```

Python

Library function identification via frequency

- Idea: identify common library functions via high usage

```
from idaapi import *  
  
file = open("function_usage.txt", "w+")  
  
functions = Functions()  
  
for f in functions:  
    name = Name(f)  
    print >> file, "%s %d" % (name, len(list(XrefsTo(f))))
```

```
sub_408FC4C8 1439  
sub_408B49C8 1476  
sub_408FC130 1502  
sub_408F8C3A 1509  
sub_408AAE3E 1654  
sub_408FC374 1680  
sub_40758D04 1715  
sub_4000A5C8 1837  
sub_408F4DFA 1935
```

Debugging

- A few different options here
 - Qcom tools like QXDM/QPST, talk to the phone over USB
 - Acquiring, licensing, ramdumps(!)
 - JTAG
 - More cost, slightly higher difficulty
 - Lauterbach Trace 32
 - Expensive, licensing, gets you as low level as you're gonna get
 - More on this later
 - Memory R/W via exploit
 - Modem Image modification

Modem image patching I

- Modem binaries are unencrypted on disk
- This facilitates easy disassembly, and easy patching
- Secboot prevents unsigned images from loading
- Signature verification performed in secure world

QSEE TrustZone Kernel Integer Overflow
Vulnerability

Dan Rosenberg

dr@azimuthsecurity.com

July 1, 2014

Modem Patching II

<https://github.com/eti1/tzexec>

Leverages a integer overflow to achieve an arbitrary write into the trustzone, and patches two bytes to neuter signature checking

Prereqs: ability to compile your own kernel and flash it

Modem internal hashes still need to be consistent

Search Results

There are **41** CVE entries that match your search.

Bits, Please!

10/08/2015

Full TrustZone exploit for MSM8974

Monday, July 24, 2017

Trust Issues: Exploiting TrustZone TEEs

Posted by Gal Beniamini, Project Zero

Bits, Please!

15/06/2016

TrustZone Kernel Privilege Escalation (CVE-2016-2431)

In this blog post we'll continue our journey from zero permissions to code execution in the TrustZone

Implementing a debugger

What are the preconditions of a debugger?

- Able to read and write from/to memory (setting breakpoints, etc)
- Breakpoints and the like implies the ability to change memory permissions
- Setting register values

How does a baseband take input?

- Over the air interface
- Shared memory
- Serial

Hayes AT commands

- Commands sent to the modem to control dialing, connection parameters, and generally manipulate things
- Extended a lot, OEM and carrier specific commands supported

```
AT
OK
AT+CGMI
Samsung Electronics
OK
```

```
AT+CLAC
AT&D
AT&F
AT&G
AT&J
AT&K
AT&L
AT&M
AT&P
AT&Q
AT&R
AT&S
AT&V
```

Implementing a debugger II

Can hook/replace AT commands

Read = AT+cmd=address,size

Write = AT+cmd=address,size,data

Just picked arbitrary commands to replace


```
AT
OK
AT+QCGQMIN=41422100,8
0x41422100=0xff

OK
AT+QCGQREQ=41422100,8,ee

OK
```

Or.....option II

TRACE32®
DOWNLOADS

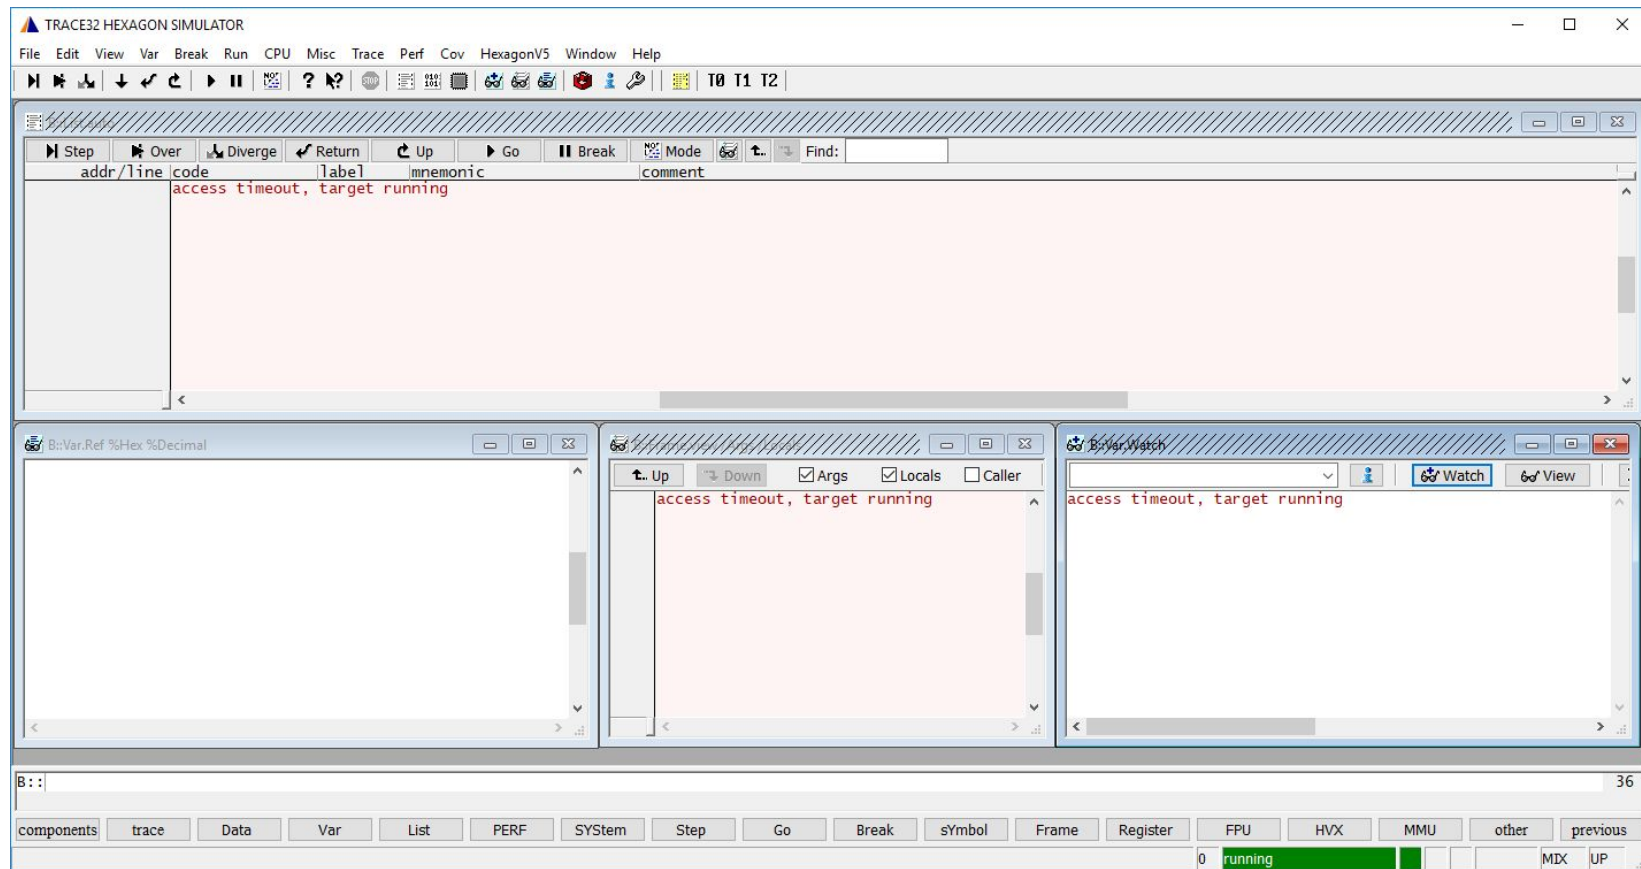


Download the latest release of the TRACE32 software, check our daily updated collection of start-up/demo scripts or get a first impression of TRACE32 by downloading an evaluation software.

- [TRACE32 Software](#)
- [TRACE32 Help System](#)
- [Start-Up and other Scripts](#)
- [Evaluation Software](#)

- Simulator for ARM/CORTEX/XSCALE
- Simulator for ARM64
- Simulator for AVR32
- Simulator for TMS320C2000
- Simulator for TMS320C5000
- Simulator for TMS320C6000
- Simulator for H8S/H8SX
- Simulator for Hexagon
- Simulator for M32R
- Simulator for M.CORE
- Simulator for MicroBlaze
- Simulator for MIPS32
- Simulator for MIPS64
- Simulator for MSP430
- Simulator for NIOS II
- Simulator for Power Architecture

Full-featured TRACE32 Instruction Set Simulators for Windows are available for free download.
Please be aware that the scripting and the remote control are limited.



Testing

- Usually need a license to broadcast on cellular frequencies (depending on country)
- Or get a Faraday cage
- Can use a Software Defined Radio (SDR) to implement our own cell stack
- Various different hardware options
 - BladeRF x40 - \$420
 - BladeRF x115 - \$650
 - USRP B200 - \$675
 - USRP x310 - \$5k

Testing II

- Quite a few open source projects have sprung up in the past few years
- YateBTS - GSM and GPRS
- OpenBTS - GSM, GPRS, 3G (UMTS)
- OpenBSC - GSM, GPRS
- OpenAirInterface - LTE
- OpenLTE - LTE
- srsLTE - LTE

Questions?

Exceptions

- Application exceptions
 - Page faults, misaligned operations, processor exceptions, etc
 - Handled by registered exception handlers
- Kernel exceptions
 - Page faults and other processor exceptions (TODO like what?)
 - Cause all execution to be terminated and the processor to be shut down
 - Processor state is saved as best it can be